

Optimal Control of a Variable-Mass Thrust-Vectoring Lander*

Nathan Ge¹

Abstract—This project studies the optimal guidance and control of a planar variable-mass rocket/lander with translational and rotational dynamics, bounded controls, nonlinear coupling, free-final-time optimal control, endpoint manifold constraints, and trajectory tracking. The project combines Pontryagin’s Maximum Principle (PMP), controllability analysis, Linear Quadratic Regulation (LQR), minimum-time optimal control, nonlinear indirect optimal control, and numerical two-point boundary value problem (TPBVP) solvers. The project is inspired by powered-descent guidance and reusable rocket landing systems studied in modern aerospace optimal control literature. All theoretical derivations must be based on PMP. Numerical implementations may use indirect shooting, multiple shooting, continuation or homotopy methods, and indirect/direct hybrid approaches, provided that the optimality conditions are derived from PMP.

github repository: [linkhttps://github.com/nzge/optimal-control-lander](https://github.com/nzge/optimal-control-lander)

I. PROBLEM FORMULATION

The state space of the lander can be represented as:

$$x = [p_x \ p_z \ v_x \ v_z \ \theta \ \omega \ m]^\top, \quad u = [\delta T \ \tau]^\top$$

Symbol	Description
<i>State Variables (x)</i>	
p_x, p_z	Horizontal and vertical position
v_x, v_z	Horizontal and vertical velocity
θ	Vehicle pitch angle
ω	Angular velocity (pitch rate)
m	Instantaneous vehicle mass
<i>Control Variables (u)</i>	
T	Commanded thrust magnitude
τ	Commanded control torque

The nonlinear equations of motion governing the lander, $\dot{x} = f(x, u)$, are defined by its kinematics, gravity, thrust resolution, and mass depletion rate.

$$\dot{x} = \begin{bmatrix} v_x \\ v_z \\ \frac{T}{m} \sin(\theta) \\ \frac{T}{m} \cos(\theta) - g \\ \omega \\ \frac{\tau}{I} \\ -\frac{\tau}{I_{sp} g_0} \end{bmatrix}$$

TABLE I: System Constraints

Mathematical Constraint	Description
<i>Control Constraints</i>	
$T \in [T_{\min}, T_{\max}]$	Actuator bounds on thrust magnitude
$\tau \in [\tau_{\min}, \tau_{\max}]$	Actuator bounds on control torque
<i>Numerically Enforced State Constraints</i>	
$p_z(t) \geq 0$	Hard deck limit preventing subsurface flight
$m(t) \geq m_{\text{dry}}$	Minimum mass limit (dry mass)

TABLE II: System Parameters

Symbol	Value	Description
g	9.81 m/s ²	Gravitational acceleration
I	10 kg · m ²	Vehicle moment of inertia
m_0	20 kg	Initial vehicle mass
α	0.005 kg/(N · s)	Propellant consumption rate
$T_{\min}$	0 N	Minimum thrust limit
$T_{\max}$	300 N	Maximum thrust limit
$\tau_{\min}$	−30 N · m	Minimum torque limit
$\tau_{\max}$	30 N · m	Maximum torque limit

II. NUMERICAL METHODOLOGY

A. Linear Systems Solver: The Sweep Method (Parts I & II)

`scipy.integrate.solve_ivp` integrates the Riccati equation backward in time to generate the feedback gain matrix $K(t)$. An interpolator was then used on this gain during the forward-pass simulation of the state trajectory.

B. Nonlinear TPBVP Solvers (Parts III & IV)

The Riccati substitution no longer applies once the optimal control problem is nonlinear, so the coupled state/costate system arising from PMP is solved directly. Part III uses *indirect single shooting* on the free-final-time mission: the $2n$ canonical equations are integrated forward and a Newton iteration adjusts the unknown initial costate, the phase-switch time, and the final time until the manifold, transversality, and $H(t_f) = 0$ residuals vanish. Part IV uses *indirect collocation* (`scipy.integrate.solve_bvp`) on the fixed-time problem: the 14 state/costate equations are discretized on a mesh and the optimal control is substituted in closed form from the stationarity conditions. Collocation is preferred for Part IV because it accepts a full state/costate *trajectory* as its initial guess, which is taken from the corresponding linearized LQR solution; this warm start lies close to the nonlinear extremal and lets the solver converge in a few mesh-refinement iterations. In both parts the actuator bounds are imposed by saturating the PMP-optimal control, consistent with Pontryagin’s pointwise minimization of H over the admissible control set.

*This work was not supported by any organization

¹Nathan Ge is a student in the Department of Mechanical and Aerospace Engineering, University of California, Los Angeles, 405 Hilgard Avenue, Los Angeles, CA 90095 nzge@g.ucla.edu

C. Constraint Enforcement Strategies

Control Constraints (T, τ): For the linear LQR phases I applied post-calculation saturation, clipping the commanded δT and τ to $[T_{\min}, T_{\max}]$ and $[\tau_{\min}, \tau_{\max}]$. For the nonlinear PMP phases the control follows in closed form from the stationarity condition $\partial H / \partial u = 0$ and is then clipped to the same actuator bounds inside the solver right-hand side; this is exactly Pontryagin's pointwise minimization of H , since for the control-affine Hamiltonian the constrained minimizer is the unconstrained stationary control projected onto the admissible box.

State Constraints ($p_z \geq 0, m \geq m_{dry}$): Adjoining state constraints directly to the PMP Hamiltonian creates severe analytical difficulties (jump discontinuities in the costates). The numerical workaround is to either:

- "Soft Constraints" (adding massive penalties to your Q matrix if the rocket goes underground)
- "Event Tracking" (terminating the ODE solver upon impact)
- Direct constraint arrays in an NLP solver.

III. PART I — FIXED-TIME LQR AROUND DIFFERENT NOMINAL TRAJECTORIES

A. Phase Formulation

Linearization: The linearized model plant is formulated as $\dot{x} = Ax + Bu$ where $A = \frac{\partial f}{\partial x}$, $B = \frac{\partial f}{\partial u}$, and $f(x, u) = \dot{x}$.

Hover

$$A_{\text{hover}} = \begin{bmatrix} 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & g & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & -\frac{g}{m_0} \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}, \quad B_{\text{hover}} = \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 0 & 0 \\ \frac{1}{m_0} & 0 \\ 0 & 0 \\ 0 & 0 \\ -\alpha & 0 \end{bmatrix}$$

Descent

$$A(t) = \begin{bmatrix} 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & g & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & -\frac{g}{m^*(t)} \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}, \quad B(t) = \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 0 & 0 \\ \frac{1}{m^*(t)} & 0 \\ 0 & 0 \\ 0 & 1/I \\ -\alpha & 0 \end{bmatrix}$$

$m^*(t) = m_0 e^{-\alpha g t}$

B. Analytical Derivations (PMP)

Hamiltonian construction

$$H(x, u, p) = L(x, u) + p^\top f(x, u) \\ \text{where } L(x, u) = x^\top Qx + u^\top Ru$$

$$H(x, u, p) = x^\top Qx + u^\top Ru + p^\top (Ax + Bu)$$

Transversality

$$p(t_f) = \left. \frac{\partial \phi}{\partial x} \right|_{t=t_f} \\ p(t_f) = 2Q_f x(t_f)$$

Adjoint Equations (costate dynamics)

$$\dot{p} = -\frac{\partial H}{\partial x}$$

Using matrix calculus (and the symmetry of Q):

$$\dot{p} = -(2Qx + A^\top p) = -2Qx - A^\top p$$

Stationarity Conditions

$$\frac{\partial H}{\partial u} = 0 \implies 2Ru + B^\top p = 0 \implies u^* = -\frac{1}{2}R^{-1}B^\top p$$

Riccati Differential equation Assumption (Riccati Substitution):

$$p(t) = P(t)x(t)$$

Where $P(t)$ is an unknown, time-varying symmetric matrix. Differentiate the assumption using the product rule:

$$\dot{p} = \dot{P}x + P\dot{x}$$

Substitute $\dot{p}$ and $\dot{x}$ (the dynamics).

$$-2Qx - A^\top Px = \dot{P}x + P(Ax + Bu^*)$$

Substitute the optimal control u^* . Remembering that $p = Px$, it becomes $u^* = -\frac{1}{2}R^{-1}B^\top Px$.

$$-2Qx - A^\top Px = \dot{P}x + P\left(Ax - \frac{1}{2}BR^{-1}B^\top Px\right)$$

$$-2Qx - A^\top Px = \dot{P}x + PAx - \frac{1}{2}PBR^{-1}B^\top Px$$

Group all terms multiplying $x(t)$. Move everything to one side of the equation:

$$\left[\dot{P} + PA + A^\top P - \frac{1}{2}PBR^{-1}B^\top P + 2Q \right] x = 0$$

Because this equation must hold true for any arbitrary state trajectory $x(t)$, the bracketed matrix expression must identically equal zero. Setting the bracket to zero yields the Riccati Differential Equation (RDE):

$$\dot{P} = -PA - A^\top P + \frac{1}{2}PBR^{-1}B^\top P - 2Q$$

Optimal State-feedback law

$$P(t_f) = 2Q_f, \quad u^*(t) = -K(t)x(t), \quad K(t) = \frac{1}{2}R^{-1}B^\top P(t)$$

C. Numerical Implementation and LQR Tuning

Q , R and Q_f matrix comparisons. We used heuristic methods to create a selection of Q , R , and Q_f that accurately reflect underlying control intuition. Fig. 1 sweeps these presets on the hover LTI plant: `pos_heavy` and `att_heavy` tighten position and attitude transients at the cost of higher $\int u^\top R u dt$; `cheap_th` and `exp_trq` isolate thrust- versus torque-dominated effort; `balanced` is the baseline used in later parts.

TABLE III: Part I.5.5 cost presets (diagonal Q , R , Q_f).

Preset	$Q = \text{diag}(\cdot)$	R	$Q_f = \text{diag}(\cdot)$
balanced	(1, 1, 0.5, 0.5, 10, 5, 0.01)	(0.1, 1)	(10, 10, 1, 1, 50, 20, 0.1)
pos_heavy	(20, 20, 2, 2, 5, 2, 0.01)	(0.1, 1)	(50, 50, 5, 5, 20, 10, 0.1)
att_heavy	(1, 1, 0.5, 0.5, 50, 25, 0.01)	(0.1, 2)	(10, 10, 1, 1, 100, 50, 0.1)
cheap_th	(1, 1, 0.5, 0.5, 10, 5, 0.01)	(0.01, 1)	(10, 10, 1, 1, 50, 20, 0.1)
exp_trq	(1, 1, 0.5, 0.5, 10, 5, 0.01)	(0.1, 10)	(10, 10, 1, 1, 50, 20, 0.1)
scaled_id	(10^{-2} , 10^{-2} , 0.04, 0.04, 25, 4, 1)	(1, 1)	same as Q




Fig. 1: Cost sweep

D. Analysis

Controllability Analysis

Algebraic Controllability Matrix (Frozen-Time): Used to check the structural control authority of the LTI system (or a "frozen" snapshot of an LTV system). For a system with $n = 7$ states:

$$C = [B \ AB \ A^2B \ A^3B \ A^4B \ A^5B \ A^6B]$$

Finite-Horizon Controllability Gramian:

$$W_c(0, t_f) = \int_0^{t_f} \Phi(0, \tau) B(\tau) B(\tau)^\top \Phi(0, \tau)^\top d\tau$$

Where $\Phi(t, \tau)$ is the state transition matrix of $A(t)$.

Numerical Computation (Differential Equation): To practically compute W_c numerically and avoid explicitly calculating the integral and state transition matrix, we can solve its differential equation form:

$$\dot{W}_c(t) = A(t)W_c(t) + W_c(t)A(t)^\top + B(t)B(t)^\top$$

Minimum Control Energy: The minimum control energy E required to reach a specific state deviation x_f is given by:

$$E = x_f^\top W_c^{-1} x_f$$




Fig. 2: Controllability Gramian and horizontal controllability proxy

Controllable energy in the weakest meaningful mode shrinks as $t_f - t$ shrinks, especially for horizontal motion (as shown by the proxy). Both trims are algebraically full rank (7/7), but the descent LTV plant at $t = 15$ s shows larger B -column entries for v_z and weaker indirect authority on p_x (horizontal acceleration enters only through attitude coupling), which is the physical origin of the "weakly controllable" label.

Controllability matrices

State order: $x = [p_x, p_z, v_x, v_z, \theta, \omega, m]^\top$

controls $u = [\delta T, \tau]^\top$

$\mathcal{C} = [B, AB, \dots, A^6B] \in \mathbb{R}^{7 \times 14}$.

Trim Snapshot $t = 0.0$ s

Descent trim: $m = 20.00$ kg, $v_x^* = 40.0$ m/s, $v_z^* = -8.0$ m/s.

TABLE IV: Controllability rank and singular values

Model	rank($\mathcal{C}$)	σ_1	σ_6	σ_7
Hover LTI	7/7	0.98	0.0489	0
Descent LTV	7/7	0.98	0.0489	0

Hover LTI at $t = 0.0$ s.

	B	AB	A^2B	A^3B	A^4B	A^5B	A^6B
p_x	0	0	0	0	0.98	0	0
p_z	0	0	0.05	0	0.00245	0	0
v_x	0	0	0	0	0.98	0	0
v_z	0.05	0	0.00245	0	0	0	0
θ	0	0	0.1	0	0	0	0
ω	0	0.1	0	0	0	0	0
m	-0.005	0	0	0	0	0	0

Descent LTV at $t = 0.0$ s.

	B	AB	A^2B	A^3B	A^4B	A^5B	A^6B
p_x	0	0	0	0	0.98	0	0
p_z	0	0	0.05	0	0.00245	0	0
v_x	0	0	0	0	0.98	0	0
v_z	0.05	0	0.00245	0	0	0	0
θ	0	0	0.1	0	0	0	0
ω	0	0.1	0	0	0	0	0
m	-0.005	0	0	0	0	0	0

Trim Snapshot $t = 15.0$ s

Descent trim: $m = 9.59$ kg, $v_x^* = 25.0$ m/s, $v_z^* = -8.0$ m/s.

TABLE V: Controllability rank and singular values

Model	rank($\mathcal{C}$)	σ_1	σ_6	σ_7
Hover LTI	7/7	0.98	0.0489	0
Descent LTV	7/7	0.98	0.1	0

Hover LTI at $t = 15.0$ s.

	B	AB	A^2B	A^3B	A^4B	A^5B	A^6B
p_x	0	0	0	0	0.98	0	0
p_z	0	0	0.05	0	0.00245	0	0
v_x	0	0	0	0	0.98	0	0
v_z	0.05	0	0.00245	0	0	0	0
θ	0	0	0.1	0	0	0	0
ω	0	0.1	0	0	0	0	0
m	-0.005	0	0	0	0	0	0

Descent LTV at $t = 15.0$ s.

	B	AB	A^2B	A^3B	A^4B	A^5B	A^6B
p_x	0	0	0	0	0.98	0	0
p_z	0	0	0.104	0	0.00511	0	0
v_x	0	0	0	0	0.98	0	0
v_z	0.104	0	0.00511	0	0	0	0
θ	0	0	0.1	0	0	0	0
ω	0	0.1	0	0	0	0	0
m	-0.005	0	0	0	0	0	0

Controllable and weakly controllable linearization comparisons.




Fig. 3: Sensitivity to initial condition error

At small perturbations: descent can be worse (weak plant + wrong linearization point). At large perturbations: hover overtakes — hover LTI may be overly optimistic at large α (linear model stays fixed while state explores larger deviations).




Fig. 4: State and control effort across trajectory




Fig. 5: Transient behavior

p_x and θ swings indicate structural difficulty. Even “controllable” hover struggles. Gramian analysis flagged as expensive to correct. The gap is clearest in p_z (mass/thrust schedule, changing trim). Descent worse




Fig. 6: Optimal solution trajectories versus time: states and control histories.




Fig. 7: Mission-plane trajectory (p_x vs p_z) with attitude vectors at sampled positions.

E. Validation




Fig. 8: Hamiltonian evolution

Because this phase utilizes the Riccati sweep method rather than a shooting solver, optimality is validated by verifying the backward-pass terminal conditions and forward-pass tracking convergence. Table VI confirms the Riccati matrix $P(t)$ strictly satisfies the terminal boundary condition $P(t_f) = Q_f$. Additionally, the closed-loop tracking error $e(t)$ is bounded within acceptable numerical tolerances at the final time.

TABLE VI: Riccati and Tracking Validation (Phase I/II)

Metric	Target	Computed	Error (Residual)
$\ P(t_f) - Q_f\ _\infty$	0	1.22×10^{-14}	1.22×10^{-14}
$\ e(t_f)\ _\infty$	0	4.51×10^{-6}	4.51×10^{-6}

IV. PART II - OPTIMAL TRAJECTORY TRACKING

A. Phase formulation

Objective (free final state, no terminal cost, as specified for Part II):

$$J = \int_0^{t_f} \left[(x - x_{\text{ref}})^T Q (x - x_{\text{ref}}) + u^T R u \right] dt$$

B. Scenario and Operating Point

The reference $x_{\text{ref}}(t)$ is one of the optimal regulation trajectories recorded in Part I (Fig. 10): a damped maneuver that settles to the hover trim. To produce a nontrivial transient, the controller is initialized at a *substantially different* state, $x_0^{\text{trk}} = [-12, 15, -4, 3, -0.25, 0.4, -0.8]^\top$, deviating from the reference start by $\Delta p_x \approx 17$ m and $\Delta \theta \approx 0.33$ rad. Because the hover linearization is altitude-invariant (p_z is decoupled and absent from every right-hand side), the maneuver is executed about a hover point $h = 12$ m above the pad; this leaves A , B , x_{ref} , and the feedback law unchanged while the absolute altitude $p_z = h + p_z^{\text{dev}}$ is used to enforce the ground constraint $p_z \geq 0$. The numerically enforced state constraints $p_z \geq 0$ and $m \geq m_{\text{dry}}$ are satisfied throughout: the minimum absolute altitude is 11.6 m (reference) and 10.0 m (tracker), and the mass margin remains above 12 kg.

C. Analytical Derivations (PMP)

The Hamiltonian:

$$H = (x - x_{\text{ref}})^\top Q(x - x_{\text{ref}}) + u^\top R u + p^\top (Ax + Bu)$$

Adjoint Equation (Costate Dynamics):

$$\dot{p} = -\frac{\partial H}{\partial x}$$

$$\dot{p} = -2Q(x - x_{\text{ref}}) - A^\top p$$

Stationarity (Optimal Control Law):

$$\frac{\partial H}{\partial u} = 0 \implies 2Ru + B^\top p = 0 \implies u^* = -\frac{1}{2}R^{-1}B^\top p$$

Transversality (Terminal Condition):

With a free final state and no terminal cost ($\phi \equiv 0$), the transversality condition gives a homogeneous terminal costate:

$$p(t_f) = \frac{\partial \phi}{\partial x} \Big|_{t=t_f} = 0$$

Tracking Riccati Equations (Riccati substitution)

The added auxiliary vector $s(t)$ accounts for the tracking offset:

$$p(t) = P(t)x(t) + s(t)$$

$$\dot{p} = \dot{P}x + P\dot{x} + \dot{s}$$

Substitute $\dot{p}$ and $\dot{x}$:

$$-2Q(x - x_{\text{ref}}) - A^\top (Px + s) = \dot{P}x + P(Ax + Bu^*) + \dot{s}$$

Substitute u ($u^* = -\frac{1}{2}R^{-1}B^\top (Px + s)$):

$$-2Qx + 2Qx_{\text{ref}} - A^\top Px - A^\top s =$$

$$\dot{P}x + PAx - \frac{1}{2}PBR^{-1}B^\top (Px + s) + \dot{s}$$

Group the terms by $x(t)$:

$$\left[\dot{P} + PA + A^\top P - \frac{1}{2}PBR^{-1}B^\top P + 2Q \right] x(t)$$

$$+ \left[\dot{s} + A^\top s - \frac{1}{2}PBR^{-1}B^\top s - 2Qx_{\text{ref}} \right] = 0$$

Since the expression must hold true for any arbitrary trajectory $x(t)$, we get the following tracking Riccati pair:

Matrix Riccati Differential Equation (The Feedback Rulebook):

$$\dot{P} = -PA - A^\top P + \frac{1}{2}PBR^{-1}B^\top P - 2Q$$

Auxiliary Vector Differential Equation (The Feedforward Tracker):

$$\dot{s} = -\left(A^\top - \frac{1}{2}PBR^{-1}B^\top \right) s + 2Qx_{\text{ref}}$$

Since $p(t_f) = P(t_f)x(t_f) + s(t_f) = 0$ must hold for an arbitrary terminal state, both pieces are integrated backward from the homogeneous terminal conditions

$$P(t_f) = 0, \quad s(t_f) = 0,$$

reflecting the absence of any terminal penalty in the Part II objective.

Closed-Loop Tracking Control Law

Once you integrate both $\dot{P}$ and $\dot{s}$ backward from t_f to 0, you can construct your final control law. Substitute $p(t) = P(t)x(t) + s(t)$ back into u :

$$u^*(t) = -\frac{1}{2}R^{-1}B^\top P(t)x(t) - \frac{1}{2}R^{-1}B^\top s(t)$$

This is elegantly split into two components:

$$u^*(t) = -K(t)x(t) + u_{ff}(t)$$

$K(t)x(t)$ is the feedback component that stabilizes the rocket and kills deviations. $u_{ff}(t)$ is the feedforward component that anticipates the control effort.

D. Analysis




Fig. 9: Tracking Performance

Tracker rejoins the reference qualitatively. Errors shrink and shapes align by mid-horizon.




Fig. 10: Recorded Part I optimal trajectory used as the reference $x_{\text{ref}}(t)$, $u_{\text{ref}}(t)$.

Quantitatively, the tracker starts with state error $\|e(0)\| \approx 19.9$, peaks at $\|e\| \approx 26.2$ near $t = 1.6$ s as it first accelerates toward the reference, and the cost-weighted error decays to within 10% of its peak by $t \approx 6.9$ s. The residual terminal error $\|e(t_f)\| \approx 1.38$ is concentrated in the vertical and mass channels ($|e_{p_z}| \approx 1.16$, $|e_{v_z}| \approx 0.67$, $|e_m| \approx 0.32$); the horizontal and attitude errors are driven to $\mathcal{O}(10^{-3})$. This residual is an expected feature of the fixed-horizon, free final-state problem with no terminal cost ($P(t_f) = 0$): as $t \rightarrow t_f$ the Riccati gain $K(t)$ and feedforward $s(t)$ both decay, so the controller “relaxes” near t_f and the weakly penalized vertical/mass states drift slightly.

The required mission-plane views (Figs. 11–12) plot p_x versus p_z with body-axis attitude arrows. The tracker spirals down from its displaced initial state and converges onto the reference at the hover target, remaining well above the ground line throughout.




Fig. 11: Reference mission trajectory in the p_x - p_z plane with attitude.




Fig. 12: Closed-loop convergence of the tracker onto the reference, with attitude arrows and the ground constraint $p_z = 0$.

Transient Response




Fig. 13: Transient Response

Scalar tracking error (cost-balanced error) decays monotonically in the weighted norm. Moving target, not zero, so no rebound mid-horizon. Position errors decay slowly and dominate; attitude error stays small throughout. Playing ‘catch-up’.




Fig. 14: Transient Control Decomposition

The control splits as $u^*(t) = -K(t)e(t) + u_{ff}(t)$. The feedback term dominates the transient and drives the large initial error to zero, while the feedforward term $u_{ff} = -R^{-1}B^T s$ is small and slowly varying because the reference is itself a decaying signal. The torque channel uses nearly the full authority at $t = 0$ ($\tau \approx 29.6$ N·m

against the ± 30 N·m bound) to arrest the attitude error, but no channel saturates over the horizon.
Control Effort




Fig. 15: Control Effort

Relative to the smooth reference plan, rejoining from the displaced initial state is expensive: the tracker expends $\int u^\top R u dt \approx 1041$ versus ≈ 86 for the reference, a $12\times$ increase concentrated in the first few seconds. Peak thrust deviation reaches $\delta T \approx 92$ N (within the $+104$ N headroom to $T_{\max}$), confirming that the chosen weights keep the catch-up maneuver feasible without thrust saturation. Once the error is nulled, the tracking control collapses onto the reference control.

Robustness to Deviations from the Nominal Trajectory




Fig. 16: Robustness to initial-condition deviations: terminal error and tracking cost as the initial offset is scaled by α .

Sweeping the initial offset by a factor α confirms the structure of the linear tracking law: the terminal error grows essentially *linearly* in α (slope ≈ 1.18), consistent with superposition of the closed-loop transition map, while the tracking cost grows *quadratically* (the cost is a quadratic form in the deviation). There is no threshold or instability over the tested range, demonstrating graceful degradation with respect to initial-condition error.

Robustness to errors in the assumed reference is examined in Fig. 17, where the reference used to build the feedforward is corrupted by zero-mean noise of level σ and the resulting controller is scored against the *true* reference. The terminal error is only mildly affected (from 1.38 at $\sigma = 0$ to 1.50 at $\sigma = 0.30$) because the feedback term rejects high-frequency reference error, but the tracking cost degrades markedly (from ≈ 3150 to

≈ 6810 , a $2.2\times$ increase) as the controller wastes effort chasing the corrupted plan.




Fig. 17: Sensitivity to a perturbed reference plan: terminal error and cost (scored against the true reference) versus reference noise level σ .

Compared with Part I regulation, the tracking problem replaces $Q_f x(t_f)$ with $P(t_f) = s(t_f) = 0$ and adds the feedforward $u_{ff}(t)$; the feedback structure is unchanged, but the controller must expend far more effort to “catch up” from a displaced initial state.

Unknown Reference Trajectory: The present law is non-causal: $s(t)$ comes from a backward integration that requires the entire future $x_{\text{ref}}(\tau)$ on $[t, t_f]$. Without that preview, only $u = -K(t)e(t)$ remains—stable, but lagging and costlier. Fig. 17 shows that reference error propagates into $s(t)$ and inflates cost even when feedback preserves terminal accuracy. A causal fix would replan online (e.g. receding-horizon MPC) or generate x_{ref} from a predictive model.

E. Validation

Because the Part II objective has no terminal cost and the final state is free, optimality is certified against the PMP necessary conditions rather than terminal convergence. The tracking Riccati pair satisfies the homogeneous free-terminal conditions $P(t_f) = 0$ and $s(t_f) = 0$, and the stationarity condition $Ru + B^\top \lambda = 0$ (with costate $\lambda = P(t)e(t) + s(t)$) holds along the entire closed-loop rollout. Table VII reports these residuals at the level of machine precision. The terminal tracking error is *not* required to vanish for this free-final-state problem; its value ($\|e(t_f)\|_\infty \approx 1.16$, dominated by the vertical/mass relaxation analyzed above) is a result, not a residual.

TABLE VII: Part II Tracking: PMP Optimality Residuals

Metric	Target	Computed	Error (Residual)
$\ P(t_f)\ _\infty$ (Riccati terminal)	0	1.0×10^{-14}	1.0×10^{-14}
$\ s(t_f)\ _\infty$ (feedforward terminal)	0	1.6×10^{-17}	1.6×10^{-17}
$\max_t \ Ru + B^\top \lambda\ $ (stationarity)	0	1.1×10^{-14}	1.1×10^{-14}

The state-equation defect of the RK4 rollout is $\mathcal{O}(10^{-1})$ at the chosen step. Figure 18 shows the Hamiltonian evolution: for the autonomous reference plan H is constant to $\sim 10^{-4}$, whereas for the tracking problem H is intentionally time-varying because the explicit dependence on $x_{\text{ref}}(t)$ makes $\partial H / \partial t \neq 0$ – the relevant optimality certificate for tracking is the stationarity/costate residual rather than constancy of H .




Fig. 18: Hamiltonian consistency: constant for the autonomous reference plan and time-varying for the tracking problem.

V. PART III - THE MISSION PROBLEM

A. Phase A — Minimum-Time Ascent

Formulation:

$$J_A = \int_0^{t_1} 1 dt = t_1$$

Analytical Derivations (PMP): Hamiltonian:

$$H_A(x, u, p) = 1 + p^T (Ax + Bu) = 1 + p^T Ax + p^T Bu$$

Adjoint equation:

$$\dot{p} = -\frac{\partial H_A}{\partial x} = -A^T p$$

Switching Function:

$$\frac{\partial H_A}{\partial u} = B^T p \implies S(t) = B^T p(t)$$

$$u_i^*(t) = \begin{cases} U_{\max,i} & \text{if } S_i(t) < 0 \\ U_{\min,i} & \text{if } S_i(t) > 0 \end{cases}$$

No singular arcs arise on thrust ($S_{\delta T}$ has one sign change). The torque channel is structurally singular for this vertical ascent ($S_\tau = p_\omega / I \equiv 0$ with $\theta = \omega = 0$), so $\tau^* = 0$ throughout.

$$H_A(t) \equiv 0 \implies 1 + p(t)^T (Ax(t) + Bu^*(t)) = 0 \quad \forall t \in [0, t_1]$$

Endpoint and Free-Final-Time Conditions: The ascent drives the vehicle to a *prescribed* rendezvous state. The first six components of the state are fixed at t_1 , while the terminal mass is left free, so the corresponding costate vanishes by transversality:

$$x_i(t_1) = x_{sky,i}, \quad i \in \{p_x, p_z, v_x, v_z, \theta, \omega\},$$

$$p_m(t_1) = 0 \quad (\text{free terminal mass}).$$

Because t_1 is free, the minimum-time transversality condition closes the system with $H_A(t_1) = 0$. The decision variables ($p(0) \in \mathbb{R}^7$, t_1) are matched against the eight residuals $\{x_{1:6}(t_1) - x_{sky,1:6}, p_m(t_1), H_A(t_1)\}$.

Implementation: Phase A is solved by indirect single shooting (mission.py). The decision vector is $Z = [p(0); t_1] \in \mathbb{R}^8$; each candidate is propagated forward with `solve_ivp` (RK45) on the 14-dimensional state-costate system using the bang-bang law above, and the eight residuals $\{x_{1:6}(t_1) - x_{sky,1:6}, p_m(t_1), H_A(t_1)\}$ are driven below 10^{-5}

by `scipy.optimize.least_squares` with 12 random multi-starts and $t_1 \in [0.3, 5]$ s. Controls enter the right-hand side as saturated PMP minimisers; $\tau^* = 0$ because S_τ is singular for the decoupled vertical ascent. The rendezvous is $x_{sky} = (0, 8, 0, v_{z,sky}, 0, 0, 0)^T$ (hover-deviation coordinates) with $v_{z,sky} \in \{-2, 0, +2\}$ m/s for the required terminal-velocity comparison.

Ascent Trajectory Results: Fig. 19 shows bang-bang thrust with exactly one switch in every case: maximum δT accelerates upward until $S_{\delta T} = (B^T p)_1$ changes sign, then minimum thrust completes the rendezvous. Thrust saturates because time appears linearly in J_A with no control penalty. Ascent time depends strongly on $v_{z,sky}$ ($t_1 = 2.39, 2.16$, and 1.98 s for $-2, 0$, and $+2$ m/s); the switch occurs late in the arc ($\approx 59\text{--}71\%$ of t_1), so deceleration to the target speed is back-loaded. Propellant use is nearly unchanged (≈ 2.08 kg) because the maneuver is short; mass depletion is too small over ~ 2 s to materially change $1/m$ and control authority.




Fig. 19: Phase A minimum-time ascent for terminal vertical velocities

$v_{z,sky} \in \{-2, 0, +2\}$ m/s: altitude $p_z(t)$, vertical velocity $v_z(t)$, bang-bang thrust deviation $\delta T(t)$, and switching function $S_{\delta T} = (B^T p)_1$. Each case exhibits a single thrust switch; the switch time, total ascent time t_1 , and propellant use vary with $v_{z,sky}$.




Fig. 20: Phase A switching function $S_{\delta T}(t)$ (top) and full mission control histories $\delta T(t)$, $\tau(t)$ (bottom) for M1 and M2; the dotted line marks the Phase A/Phase B handoff t_1 . The single sign change of $S_{\delta T}$ confirms the bang–bang structure with one switch.

B. Phase B — Return-and-Land Problem

Phase B begins precisely where Phase A ends, thus the initial condition for the landing phase is $x(0) = x_1$.

Formulation & Core PMP:

$$J_B = \int_0^{t_f} (x^\top Qx + u^\top Ru) dt$$

Generic Hamiltonian, and the fundamental $H(t_f) = 0$ condition:

$$H_B(x, u, p) = x^\top Qx + u^\top Ru + p^\top (Ax + Bu)$$

The unconstrained control law and adjoint equations remain identical to Phase I:

$$u^*(t) = -\frac{1}{2}R^{-1}B^\top p(t)$$

$$\dot{p}(t) = -2Qx(t) - A^\top p(t)$$

Because t_f is free, PMP imposes the free-time transversality condition. The Hamiltonian evaluated at the terminal time must equal zero:

$$H_B(t_f) = 0 \implies x(t_f)^\top Qx(t_f) + u(t_f)^\top Ru(t_f) + p(t_f)^\top (Ax(t_f) + Bu(t_f)) = 0$$

Transversality:

$$p(t_f) = \frac{\partial \phi}{\partial x} \Big|_{t_f} + \sum_j v_j \nabla h_j(x(t_f))$$

Manifold 1 (Pinpoint Landing)

Description

$$h_1(x) = p_z = 0$$

$$h_2(x) = v_z = 0$$

$$h_3(x) = \theta = 0$$

Gradients

$$\nabla h_1 = [0, 1, 0, 0, 0, 0]^\top$$

$$\nabla h_2 = [0, 0, 0, 1, 0, 0]^\top$$

$$\nabla h_3 = [0, 0, 0, 0, 1, 0]^\top$$

Transversality

$$p_{p_x}(t_f) = \frac{\partial \phi}{\partial p_x} \quad (\text{Free state})$$

$$p_{p_z}(t_f) = \frac{\partial \phi}{\partial p_z} + v_1 \quad (\text{Constrained})$$

$$p_{v_x}(t_f) = \frac{\partial \phi}{\partial v_x} \quad (\text{Free state})$$

$$p_{v_z}(t_f) = \frac{\partial \phi}{\partial v_z} + v_2 \quad (\text{Constrained})$$

$$p_\theta(t_f) = \frac{\partial \phi}{\partial \theta} + v_3 \quad (\text{Constrained})$$

$$p_\omega(t_f) = \frac{\partial \phi}{\partial \omega} \quad (\text{Free state})$$

$$p_m(t_f) = \frac{\partial \phi}{\partial m} \quad (\text{Free state})$$

Manifold 2 (Surface Landing)

Description: landing on the surface of a dome/sphere.

$$h_1(x) = (p_x - p_c)^2 + p_z^2 - r^2 = 0$$

$$h_2(x) = v_z = 0$$

$$h_3(x) = \theta = 0$$

The Gradients: h_1 couples p_x and p_z , spanning 2 dimensions.

$$\nabla h_1 = [2(p_x - p_c), 2p_z, 0, 0, 0, 0]^\top$$

$$\nabla h_2 = [0, 0, 0, 1, 0, 0]^\top$$

$$\nabla h_3 = [0, 0, 0, 0, 1, 0]^\top$$

Transversality

$$p_{p_x}(t_f) = \frac{\partial \phi}{\partial p_x} + 2v_1(p_x(t_f) - p_c)$$

$$p_{p_z}(t_f) = \frac{\partial \phi}{\partial p_z} + 2v_1 p_z(t_f)$$

(The remaining 5 costates are identical to M1)

Conclusion for M2: We can eliminate the unknown multiplier v_1 by dividing the first two equations (assuming $\frac{\partial \phi}{\partial x} = 0$). This yields the exact boundary condition passed to the TPBVP solver:

$$\frac{p_{p_x}(t_f)}{p_{p_z}(t_f)} = \frac{p_x(t_f) - p_c}{p_z(t_f)}$$

Numerical Implementation: The free-final-time landing problem is solved as a true indirect PMP boundary value problem by collocation. The combined state–costate system ($\dot{x} = Ax + Bu^*$, $\dot{p} = -2Qx - A^\top p$) with the stationarity law $u^* = \text{sat}(-\frac{1}{2}R^{-1}B^\top p)$ is mapped onto $s \in [0, 1]$ via $t = t_1 + s(t_f - t_1)$, and the unknown duration $t_f - t_1$ enters as a free parameter. The split boundary conditions are the handoff $x(t_1) = x_1$, the seven manifold endpoint and transversality residuals (Section 7.3: $\{p_z, v_z, \theta, p_{p_x}, p_{v_x}, p_\omega, p_m\}$ for M1; $\{h_1, v_z, \theta, p_{v_x}, p_\omega, p_m, \text{collinearity}\}$ for M2), and the free-time closure $H_B(t_f) = 0$, giving the fifteen conditions required for the fourteen states plus the parameter. The platform is centred at $(p_c, 0) = (0, 0)$ so that M1 (flat ground) lands at $(0, 0)$ and M2 (circular platform) lands at the apex $(0, r)$; both share the same horizontal footprint, isolating the effect of the terminal-manifold geometry. Among the multiple t_f local optima satisfying the necessary conditions, the global minimiser of J_B that respects $p_z \geq 0$ throughout is retained.




Fig. 21: Mission-plane trajectories p_x vs p_z with body-attitude vectors at sampled positions: M1 flat-ground landing (left, touchdown at the origin) and M2 circular-platform landing (right, touchdown at the platform apex). The ascent (Phase A) reaches x_{sky} (star); the descent (Phase B) tracks the optimal landing arc.




Fig. 22: Overlaid M1 and M2 mission planes. Both scenarios share the launch site, x_{sky} , and horizontal footprint; they differ only in the terminal manifold (flat ground at $(0,0)$ vs. circular platform apex at $(0,r)$).

Landing Manifold Comparison: Both missions share Phase A, x_{sky} , and the horizontal footprint; only the Phase B terminal manifold differs (Fig. 22). Fig. 23 and Table VIII summarise the five comparison axes:

Landing time. M2 completes the mission sooner: $t_f = 7.32$ s versus 8.10 s for M1 ($\Delta t_f \approx 0.78$ s, Table VIII). Phase B alone is also shorter (5.16 vs. 5.94 s) because touchdown at the platform apex $(0,r) = (0,5)$ m avoids the long final vertical flare required to reach $p_z = 0$ on flat ground.

Energy consumption. Despite the shorter clock time, M2 costs more control energy: $\int_{t_1}^{t_f} u^\top R u dt \approx 526$ versus 216 for M1 ($\approx 2.4\times$, Table VIII). The lofted descent arc and stronger horizontal-vertical coordination dominate the energy budget.

Switching structure. Phase A is identical in both cases (one bang-bang thrust switch at $t_1 = 2.16$ s). Phase B uses smooth saturated costate feedback in both manifolds—no further bang-bang arcs. Fig. 20 shows similar overall control profiles, with M2 displaying larger mid-descent δT and τ activity to steer onto the curved approach.

Terminal sensitivity. Perturbing the Phase A $\rightarrow$ B handoff state by ± 0.05 on each component (holding the optimal costate and t_f fixed) raises the worst-case terminal residual to ~ 51 for M1 but $\sim 1.9 \times 10^3$ for M2. The sphere constraint and collinearity condition $p_{p_x}(t_f) p_z(t_f) = p_{p_z}(t_f) (p_x(t_f) - p_c)$ tightly couple position and costate at touchdown on M2.

Feasibility. Both scenarios satisfy $p_z \geq 0$ along the trajectory, converge with $\|res\|_\infty \sim 6.2 \times 10^{-8}$ (below the 10^{-5} target), and meet the terminal manifold and $H_B(t_f) = 0$ residuals in Table VIII. Touchdown occurs at $(0,0)$ and $(0,5)$ m as intended.

Effect of manifold geometry. Flat ground $M_1 = \{p_z = 0, v_z = 0, \theta = 0\}$ allows a nearly vertical descent over the shared footprint (Fig. 21, left). The platform manifold $M_2 = \{h_1(x) = 0, v_z = 0, \theta = 0\}$ replaces the altitude fix with the sphere $h_1 = (p_x - p_c)^2 + p_z^2 - r^2 = 0$, forcing a bowed trajectory that intercepts the surface along its outward normal at the apex; transversality requires the position costate to align with that normal. The resulting arc demands sustained lateral thrust and attitude correction through Phase B (Fig. 21, right), which—together with the higher Q weights used for M2—explains the shorter but energetically costlier landing relative to M1.




Fig. 23: M1 vs M2 comparison: total landing time t_f , Phase B duration, Phase B control energy $\int u^\top R u dt$, Phase A thrust switches, shooting feasibility $\|res\|_\infty$, and terminal sensitivity to ± 0.05 handoff perturbations.

Validation: All Part III necessary conditions are verified to numerical tolerance: the terminal manifold and transversality residuals, the free-final-time condition $H_B(t_f) = 0$, the Hamiltonian consistency $H_A \equiv 0$ (minimum time) and $H_B \equiv 0$ (free final time, autonomous), and the state constraint $p_z \geq 0$ along the full trajectory. Table VIII reports the computed residuals.

TABLE VIII: Part III Mission: PMP / TPBVP Validation Residuals

Metric	M1 (flat)	M2 (platform)
Shooting $\ res\ _\infty$ (target 10^{-5})	6.2×10^{-8}	6.2×10^{-8}
Terminal manifold $\ res\ _\infty$	7.0×10^{-21}	1.1×10^{-14}
Free-time condition $H_B(t_f)$	-7.0×10^{-21}	0.0
H_A drift (min-time, $\equiv 0$)	2.9×10^{-9}	2.9×10^{-9}
H_B drift (free-time, $\equiv 0$)	9.7×10^{-12}	9.6×10^{-11}
Min. altitude p_z [m] (≥ 0)	0.000	0.000
Ascent time t_1 [s]	2.158	2.158
Landing time t_f [s]	8.097	7.315
Phase B energy $\int u^\top R u dt$	215.7	526.2
Touchdown (p_x, p_z) [m]	(0,0)	(0,5)




Fig. 24: Hamiltonian consistency: $H_A \equiv 0$ on Phase A (minimum time) and $H_B \equiv 0$ on Phase B (free final time), both to numerical tolerance.

VI. PART IV - THE NONLINEAR OPTIMAL CONTROL NEAR HOVER

A. Phase Formulation

Consider the fixed-time, free-endpoint nonlinear optimal control problem:

$$\min_u \int_0^{t_f} (x^\top Q x + u^\top R u) dt$$

Recall: State Equations (which is also $\dot{x} = \frac{\partial H}{\partial p}$):

$$\dot{x} = \begin{bmatrix} v_x \\ v_z \\ -\frac{T}{m} \sin \theta \\ \frac{T}{m} \cos \theta - g \\ \omega \\ \frac{\tau}{I} \\ -\alpha T \end{bmatrix}$$

System is anchored by the known initial launch state at $t = 0$:

$$x(0) = x_0$$

B. Analytical Derivations

Nonlinear Hamiltonian

Adjoin the nonlinear dynamics to the running cost using the costate vector $p = [p_1, p_2, p_3, p_4, p_5, p_6, p_7]^\top$:

$$H(x, u, p) = x^\top Q x + u^\top R u + p_1 v_x + p_2 v_z - p_3 \frac{T}{m} \sin \theta + p_4 \left(\frac{T}{m} \cos \theta - g \right) + p_5 \omega + p_6 \frac{\tau}{I} - p_7 \alpha T$$

Adjoint equations (costate equations)

$$\dot{p} = -\frac{\partial H}{\partial x}$$

$$\dot{p}_1 = -2Q_{11}p_x$$

$$\dot{p}_2 = -2Q_{22}p_z$$

$$\dot{p}_3 = -2Q_{33}v_x - p_1$$

$$\dot{p}_4 = -2Q_{44}v_z - p_2$$

$$\dot{p}_5 = -2Q_{55}\theta + p_3 \frac{T}{m} \cos \theta + p_4 \frac{T}{m} \sin \theta$$

$$\dot{p}_6 = -2Q_{66}\omega - p_5$$

$$\dot{p}_7 = -2Q_{77}m - p_3 \frac{T}{m^2} \sin \theta + p_4 \frac{T}{m^2} \cos \theta$$

Stationarity conditions ($\frac{\partial H}{\partial u} = 0$)

Consistent with the implementation and with Parts I–II, the running cost penalizes the *deviation* control $u = [\delta T, \tau]^\top$, where $\delta T = T - m_0 g$ is the thrust increment about the hover trim and the thrust entering the equations of motion is $T = m_0 g + \delta T$ (so $\partial T / \partial \delta T = 1$). Differentiating H directly with respect to the deviation control yields the optimal increment, with no trim offset subtracted afterwards:

For the thrust increment (δT):

$$\frac{\partial H}{\partial \delta T} = 2R_{11} \delta T - p_3 \frac{\sin \theta}{m} + p_4 \frac{\cos \theta}{m} - p_7 \alpha = 0$$

$$\delta T^* = \text{sat}_{\delta T_{\min}}^{\delta T_{\max}} \left[\frac{1}{2R_{11}} \left(p_3 \frac{\sin \theta}{m} - p_4 \frac{\cos \theta}{m} + p_7 \alpha \right) \right],$$

where the saturation limits are the actuator bounds expressed in deviation form, $\delta T \in [T_{\min} - m_0 g, T_{\max} - m_0 g]$.

For Torque (τ):

$$\frac{\partial H}{\partial \tau} = 2R_{22} \tau + \frac{p_6}{I} = 0$$

$$\tau^* = \text{sat}_{\tau_{\min}}^{\tau_{\max}} \left[-\frac{p_6}{2IR_{22}} \right]$$

Transversality (terminal condition)

Because the terminal state is subject to optimization, the terminal costate at $t = t_f$ is governed by the transversality condition based on the terminal penalty $\phi(x(t_f))$ (or specific manifold constraints):

$$p(t_f) = \frac{\partial \phi}{\partial x} \Big|_{t=t_f}$$

Complete TPBVP

The complete TPBVP is formed by substituting the optimal, saturated control laws (δT^* , τ^*) back into the 7 state equations and 7 adjoint equations. This creates a deeply coupled system of 14 nonlinear, first-order ordinary differential equations. This system must be solved simultaneously across the domain $t \in [0, t_f]$, subject to the split boundary conditions $x(0)$ and $p(t_f)$.

$$\begin{aligned}
\dot{p}_x &= v_x \\
\dot{p}_z &= v_z \\
\dot{v}_x &= -\frac{T}{m} \sin \theta \\
\dot{v}_z &= \frac{T}{m} \cos \theta - g \\
\dot{\theta} &= \omega \\
\dot{\omega} &= \frac{\tau}{I} \\
\dot{m} &= -\alpha T
\end{aligned} \tag{1}$$

$$\begin{aligned}
\dot{p}_1 &= -2Q_{11}p_x \\
\dot{p}_2 &= -2Q_{22}p_z \\
\dot{p}_3 &= -2Q_{33}v_x - p_1 \\
\dot{p}_4 &= -2Q_{44}v_z - p_2 \\
\dot{p}_5 &= -2Q_{55}\theta + \frac{T}{m}(p_3 \cos \theta + p_4 \sin \theta) \\
\dot{p}_6 &= -2Q_{66}\omega - p_5 \\
\dot{p}_7 &= -2Q_{77}m + \frac{T}{m^2}(p_4 \cos \theta - p_3 \sin \theta)
\end{aligned}$$

C. Numerical Implementation

The Part IV problem is solved in hover-deviation coordinates (x is the deviation from the hover trim $T^* = m_0 g$, $\theta^* = 0$; control $u = [\delta T, \tau]^\top$), so that the linearized solutions of the previous parts can serve directly as the nonlinear warm start. The implementation (`nonlinear.py`) builds three comparison scenarios on a common initial condition x_0 and a common short horizon:

Restriction to the near-hover regime. Because true hover is *not* an equilibrium of a variable-mass vehicle (holding $T = mg$ continuously depletes mass at $\dot{m} = -\alpha mg$), the horizon is kept short, $t_f = 4$ s with 201 uniform nodes. This keeps the trajectory near hover and the mass safely above m_{dry} , which is the validity regime prescribed for this part.

Scenario 1 — LQR on the linear plant. The backward Riccati sweep of Part I produces $P(t)$ and the feedback $u = -R^{-1}B^\top P(t)x$, integrated forward on $\dot{x} = Ax + Bu$.

Scenario 2 — LQR on the nonlinear plant. The *same* time-varying gain is applied to the full nonlinear EOM, integrated with a fixed-step RK4 scheme, with control saturation enforced by clipping. This isolates the modeling error incurred by using the linear law on the true plant.

Scenario 3 — Nonlinear PMP / TPBVP. The coupled 14-state state/costate system of Eq. (1) is solved by collocation (`scipy.integrate.solve_bvp`). The optimal control is substituted in closed form from the stationarity conditions and saturated to the actuator bounds. The split boundary conditions are the initial-state match $x(0) = x_0$ and the free-endpoint transversality $p(t_f) = 2Q_f x(t_f)$.

Warm start and costate scaling. The collocation guess uses the Scenario 1 state history and the costate guess $p(t) = 2P(t)x(t)$; the factor of two maps the $\frac{1}{2}$ -scaled Riccati costate to the no- $\frac{1}{2}$ running cost $x^\top Qx + u^\top Ru$

used here. Convergence is declared at a collocation tolerance of 10^{-6} .

D. Analysis

Three scenarios share $x_0 = [1, 1.5, 0.4, -0.25, 0.12, 0, 0.05]^\top$ and $t_f = 4$ s: (1) LQR on the LTI plant, (2) the same LQR law on the nonlinear plant, and (3) nonlinear PMP/collocation on the nonlinear plant with an LQR warm start.




Fig. 25: State trajectories for the three scenarios.




Fig. 26: Mission plane p_x vs. p_z with attitude vectors.




Fig. 27: Thrust and torque histories.




Fig. 28: Control effort $J_u = \int u^T R u dt$ and total cost J .




Fig. 29: Terminal state norm and components at t_f .




Fig. 30: Collocation convergence and tolerance sensitivity.




Fig. 31: Control effort and PMP savings vs. scaled initial pitch θ_0 .




Fig. 32: Per-component TPBVP boundary and optimality residuals.

Figs. 25–26 show scenarios (1) and (2) agree early but diverge as θ and m evolve; the nonlinear PMP solution

trims peak p_z and θ by exploiting $\sin \theta$ and T/m coupling. The linear law over-commands on the true plant (Fig. 27), raising J_u from 13.2 to 42.9 while PMP recovers 39.4 ($\approx 8\%$ savings) and lowers total J by 41% vs. scenario (2) (Fig. 28). Terminal norms tighten under PMP (Fig. 29). Collocation converges in few refinements from the LQR seed (Fig. 30); the PMP advantage grows with initial pitch (Fig. 31). Boundary residuals meet the 10^{-6} target (Fig. 32). The linear approximation remains accurate for small attitude/mass deviations over 4 s; it breaks down once thrust-gravity resolution and mass depletion accumulate.

E. Validation

To verify the numerical solution to the Two-Point Boundary Value Problem (TPBVP) and prove strict adherence to Pontryagin’s Maximum Principle (PMP), the solution was validated against the necessary optimality conditions.

Hamiltonian Consistency: Part IV is a *fixed-final-time*, free-endpoint problem, and the hover-trim dynamics are autonomous. The Hamiltonian must therefore remain *constant* along the optimal trajectory; unlike the free-final-time mission of Part III, this constant is *not* required to vanish. As shown in Figure 33, the Hamiltonian of the nonlinear PMP solution is constant to a drift of $\mathcal{O}(10^{-9})$ about $H(t_f) \approx 34.7$, confirming Hamiltonian consistency. By contrast, applying the linear LQR law to the nonlinear plant produces a Hamiltonian that drifts by $\mathcal{O}(10^2)$ over the same horizon—a direct, quantitative certificate that the linear law is *not* an extremal of the nonlinear problem.




Fig. 33: Hamiltonian evolution $H(t)$ for the three scenarios. The nonlinear PMP solution (bottom) holds H constant to $\mathcal{O}(10^{-9})$ as required for an autonomous fixed-time extremal, whereas the LQR law on the nonlinear plant drifts by $\mathcal{O}(10^2)$.

Terminal Conditions, Adjoint Boundaries, and Shooting Residuals: Because the terminal state is free, the boundary conditions verified are the initial-state match $x(0) = x_0$ and the free-endpoint transversality $p(t_f) = 2Q_f x(t_f)$; there are no fixed terminal-state targets in this part. Table IX reports the per-component boundary residuals together with the interior collocation defect and the Hamiltonian drift.

The global $\mathcal{L}_\infty$ norm of the boundary-condition residual was 1.8×10^{-15} and the maximum interior collocation residual was 9.7×10^{-7} on a converged mesh of 219 nodes—both within the requested solver tolerance, confirming that the computed solution satisfies the necessary optimality conditions to numerical precision.

TABLE IX: Part IV validation: boundary and optimality residuals of the nonlinear TPBVP (fixed-time, free endpoint).

Type	Boundary / optimality condition	Residual
Initial	$x(0) = x_0$ ($\mathcal{L}_\infty$)	3.9×10^{-26}
Transv.	$p_{p_z}(t_f) = 2Q_{f,22} p_z(t_f)$	1.8×10^{-15}
Transv.	$p_\theta(t_f) = 2Q_{f,55} \theta(t_f)$	8.9×10^{-16}
Transv.	$p_\omega(t_f) = 2Q_{f,66} \omega(t_f)$	5.6×10^{-17}
Transv.	$p(t_f) = 2Q_f x(t_f)$ ($\mathcal{L}_\infty$)	1.8×10^{-15}
Hamiltonian	$H(t)$ drift (autonomous, fixed t_f)	1.4×10^{-9}
Collocation	interior RMS defect (219 nodes)	9.7×10^{-7}

ACKNOWLEDGMENT

The author thanks Professor Shahriar Talebi for guidance and instruction throughout the MAE 270C Optimal Control course. Their insights into optimal control theory formed the foundation of this project.

REFERENCES

- [1] Liberzon, D. (2012). *Calculus of variations and optimal control theory: A concise introduction*. Princeton University Press.
- [2] Bryson, A. E., & Ho, Y.-C. (2018). *Applied optimal control*. Routledge. <https://doi.org/10.1201/9781315137667>
- [3] Kirk, D. E. (2004). *Optimal control theory: An introduction*. Courier Corporation.
- [4] Acikmese, B., & Ploen, S. R. (2007). Convex programming approach to powered descent guidance for Mars landing. *Journal of Guidance, Control, and Dynamics*, 30(5), 1353–1366. <https://doi.org/10.2514/1.27553>